

Maui-as-Product: A Bus-Mediated Architecture for Composing Agentic Capabilities

Status: Design white paper · v1 · 2026-05-29

Author: SPY (Sean Young) + Cosmo

Scope: How to turn a private agentic colleague (Maui) and an orchestration framework (Mycelium) into a distributable product, by composing a large existing capability fleet over a single message bus — without losing CLI-level control and without letting the system become a junk drawer.

1. Executive summary

We have three assets that today do not know about each other:

1. **Mycelium** — a language-agnostic, patent-pending orchestration framework. It knows how to decompose a goal into many parallel agents (*leaves*) and compose their output. This is the reusable orchestration grammar.
2. **The Frank/Maui capability fleet** — ~789 registered capabilities plus a large RAG/research-archive layer. These know how to *do* specific things. They are capabilities, not orchestration.
3. **Maui's front doors** — iMessage and phone today; web/email/API later. These know how to *receive intent*. They are transport.

The thesis of this paper: **a single message bus is the seam that lets all three compose without coupling them**. If every capability, every RAG, every front door, and every orchestration run speaks to the same peer-to-peer bus using one protocol, then integration stops being an N-by-N problem.

Be precise about what collapses and what does not. Integration has two parts: **transport** (how you reach a thing — HTTP vs MCP vs RAG-query) and **interface** (where it lives, what request shape it expects, what it returns). The bus collapses the *transport* cost to a small, fixed set of **transport adapters** — roughly four (HTTP service, RAG/archive, MCP, registry). It does **not** make the *interface* cost vanish: each capability's endpoint, request schema, and result shape still has to be described somewhere. That description lives as a **per-capability interface descriptor in cap-reg** (the registry), authored once per capability — declaratively, not as bespoke integration code.

So the honest claim is: **~4 transport adapters + one declarative interface descriptor per capability in cap-reg**. The win is real — interface facts become uniform, declarative registry data instead of N hand-written integration modules, and any capability that already has a descriptor is callable for free — but the per-capability cost moves into cap-reg rather than disappearing. Any implementation plan must budget for authoring/maintaining those descriptors.

Two principles keep the system honest:

- **Mycelium is the orchestration engine, above the bus — invokable two ways.** A cultivation is an orchestrator that *uses* the bus as a client and is **not** automatically a biome. What is load-bearing is the *engine* (deterministic fan-out, CommitQueue, dry-run, spend-cap, audit trail), **not** a human at a terminal. The same engine is invoked whether the operator runs `mycelium ...` by hand or a routed request triggers it programmatically. The CLI is one front door onto the engine; the product is another. The gates (spend-gate, promotion) are what make auto-triggered, headless invocation safe.
- **The bus stays curated.** Cultivations *produce candidates* for new permanent capabilities, but joining the organism requires an explicit **promotion gate** (manual to start, policy-driven later). Nonsense a user builds never auto-joins.

The first product surface is iMessage + phone. That is just the first *front-door biome* — not a special case. Everything else in this paper is the machinery that makes adding the second and third front door cost almost nothing.

2. Terminology

This paper uses Mycelium's biological vocabulary. Names are not load-bearing — read the seams, not the labels — but here is the working glossary.

Term	Meaning
Nutrient	A unit of work or data that flows on the bus. In practice: an <i>intent</i> (what a user asked) travelling in, or a <i>result</i> travelling back.
Biome	Any participant plugged into the bus. A front door, a capability, a RAG — all biomes. Mental model: "a node that can send and receive nutrients."
The Bus	A peer-to-peer message fabric. No central boss. Every biome speaks to it in one shared protocol. (Today: BIOME BUS, port 9450. The bus concept is integrated later in the build sequence; this paper designs forward-compatible.)
Capability	Something that <i>does</i> work: a Frank service, a RAG, an MCP tool. ~789 exist already.
Transport adapter	A thin shim that makes an entire <i>class</i> of capabilities (all HTTP services, all RAGs, all MCP tools) reachable over the bus. Written once per protocol , not once per capability. Handles <i>transport</i> only — see <i>Interface descriptor</i> . ~4 of them.
Interface descriptor	Per-capability registry data (endpoint, request schema, result shape) held in cap-reg. Authored once per capability , declaratively. The transport adapter + the descriptor together make a capability callable. The reusability is "~4 adapters + 1 descriptor per capability" — transport cost is fixed; interface cost is real but declarative (see §1, §3.4).
Cultivation	A Mycelium run that decomposes a <i>novel</i> goal from scratch into many parallel agents (leaves) and composes their output. The "build-anything" path — Tier 3 .
Composed task	Tier 2. A request needing several capabilities combined in a <i>known</i> shape — handled by a pre-defined, contract-bound Mycelium agent (no goal decomposition). The "middle weight" between a single lookup and a full cultivation.

Term	Meaning
Triage	The sizing-and-routing step. Runs once, up front. Picks a tier — Tier 1 lookup (single) · Tier 2 composed task · Tier 3 cultivation — and routes to it. Triage never composes ; composition is always owned downstream by a contract-bound agent or a cultivation.
Promotion	The <i>gated</i> path by which a proven cultivation output becomes a permanent, reusable capability-biome. Never automatic.

3. Structural view — the organism

The structural view answers: *where do things live, and how do they relate?*

Maui-as-Product — v3 (terminology key + curated promotion)

Resolved: cultivations run privately by default; becoming a reusable biome requires an approval gate. The bus stays curated.

TERMINOLOGY — READ THIS FIRST

Nutrient — a unit of work/data that flows on the bus. Here: an *intent* (what the user asked) or a *result* coming back.

The Bus — peer-to-peer message fabric. No central boss. Every biome talks to it in one shared protocol.

Adapter — thin shim that makes a whole CLASS of capabilities (all HTTP services, all RAGs...) speak bus. **Written once per protocol, not per capability.**

Triage — the sizing step: is this a cheap lookup, a single capability, or a full cultivation? Decided once, up front.

Biome — any participant plugged into the bus. A front door, a capability, a RAG — all biomes. Think: "a node that can send/receive nutrients."

Capability — something that DOES work: a Frank service, a RAG, an MCP tool. ~789 of them exist already.

Cultivation — a Mycelium run that decomposes a big goal into many parallel agents (*leaves*) and composes their output. The "build-anything" path.

Promotion — the GATED path by which a proven cultivation output becomes a permanent, reusable capability-biome. Not automatic.

1 · FRONT-DOOR BIOMES — RECEIVE INTENT + IDENTITY

iMessage / phone v1

Injects an intent nutrient with reply-channel + caller identity attached.

CC session operator

Your door. Same injection, but skips the spend-gate — you are the gate.

web / email / API later

Each new channel = one more front-door biome. Core untouched.

▼ intent

2 · TRIAGE — SIZE & ROUTE (NEVER COMPOSES)

⊙ Lookup ≈ \$0

Single capability / model call / RAG query. Answer inline, done.

⊙ Composed task

2+ capabilities, KNOWN shape → routes to a **contract-bound Mycelium agent**. Pre-defined first, dynamic later. Agent composes, not triage.

⊙ Build → spend-gate

Novel goal, decompose from scratch → cultivation. Over threshold → texts "build X, ~\$Y, go?", waits.

Triage picks the tier and hands off. Composition is ALWAYS owned downstream (agent or cultivation), never by triage.

▼ routed + approved work

⚡ THE BUS — peer-to-peer · no central node · one protocol
The reusable seam. Everything on this page speaks to it the same way.

{ intent, origin_channel, identity, budget_cap, trace_id, deadline, reply_to }

▲▼ request / response

3 · CAPABILITY BIOMES — 789 + RAGS VIA ~4 ADAPTERS

HTTP-service

Frank services (research, recon, comp-intel...).

RAG / archive

research-archive, knowledge-router, topic RAGs.

MCP

Gmail, Drive, Brave, Shodan, any MCP tool.

cap-reg

The registry/menu. Triage queries it; biomes self-announce.

~4 TRANSPORT adapters + 1 interface descriptor per capability (in cap-reg). Transport cost is fixed; interface cost is declarative, not eliminated.

▼ the "build-anything" path lands here · leaves reach back UP to capability biomes

4 · MYCELIUM ENGINE — OWNS ALL COMPOSITION · INVOKED BY HAND OR AUTO-TRIGGERED

⊙ Tier-2 contract-bound agent

Fixed-shape composition via NUTRIENTS contract. No decomposition. Cheap, predictable, testable. v1 = pre-defined; dynamic later.

⊙ Tier-3 cultivation

plant→freeze→cultivate→harvest. Novel goal decomposed from scratch. Owns fan-out/CommitQueue. NOT automatically a biome.

🔥 Leaves/steps call capability-biomes

Toolset = capability-biomes on the bus. Where composition + the fleet finally meet.

CLI is the engine's interface, NOT a human-at-terminal requirement. Same build cycle whether you type "mycelium ..." or a routed request auto-triggers it. Spend-gate + promotion gate make headless invocation safe.

Resolved: not every cultivation joins the organism. A user could build nonsense — it stays private until it earns promotion. ↓

▼ only IF it proves useful

5 · PROMOTION — GATED PATH FROM CULTIVATION → PERMANENT BIOME

🔥 Cultivation output

A working fruit. Lives privately. Reusable only by re-running.

🔒 Approval gate

Operator (or policy) reviews: useful? safe? non-nonsense? Worth keeping?

🔧 Wrapped as adapter

Gets a bus adapter + cap-reg entry. Now discoverable.

🌱 Permanent capability-biome

Joins layer 3. Future triage can route to it directly — no rebuild.

The bus stays **CURATED**. Cultivations produce candidates; only approved ones become permanent. Nonsense never auto-joins.

CROSS-CUTTING — SPANS EVERY LAYER

🧠 Memory / identity

Who's asking? What does Maui know about them + the org? (colleague+product boundary lives here.)

💰 Budget ledger

Per-trace cap from the envelope. Triage reads it; cultivations decrement it.

🔍 Observability

trace_id threads the whole journey: door → triage → bus → biomes → reply.

🛡️ Safety tier

Product caller ≠ operator. Hides sensitive biomes (OSINT/recon) from the product tier.

(Diagrams are reproduced in §3.1–§3.6 as text so the paper is self-contained without the rendered image.)

3.1 Layer 1 — Front-door biomes (receive intent + identity)

Front doors receive intent and nothing more. They do **not** decide what to do. Each one wraps an inbound message as an intent nutrient and attaches:

- a **reply channel** (so the answer knows how to get home),
- a **caller identity** (so memory and the safety tier know who is asking).
- **iMessage / phone** — v1. The first product surface.
- **CC session** — the operator door. Same injection, but it *skips the spend-gate* because the operator is the gate.
- **web / email / API** — later. Each new channel is one more front-door biome. The core is untouched when you add one.

3.2 Layer 2 — Triage (classify depth, route — never compose)

Triage **sizes and routes**, and it runs **once** per request. Critically, triage does **not** itself orchestrate or compose multiple capability calls — composition is always owned by something downstream (a contract-bound agent or a cultivation). Triage's only job is to pick the right tier and hand off. It sizes a request into one of **three tiers**:

- **Tier 1 — Lookup (≈ \$0)**. One model call or one RAG query, or a dispatch to a *single* capability-biome. Triage routes it and the answer comes straight back. No composition.
- **Tier 2 — Composed task**. A request that needs *several* capability-biomes combined (sequence or parallel) but is **not** a novel build. Triage routes it to a **contract-bound Mycelium agent** — a fixed-shape composition (call these capabilities, compose this way) defined ahead of time with a NUTRIENTS contract. The agent owns the orchestration; triage just selected it. (See §3.5 for why this is Mycelium, not triage, doing the work; and §6.2 for the pre-defined→dynamic evolution.)
- **Tier 3 — Build-anything**. A novel goal that must be decomposed from scratch → a full Mycelium cultivation. Expensive; gated (see §3.7 and §6).

This three-tier split resolves the "middle weight" directly: **not every multi-capability request is a cultivation**. A composed task (Tier 2) is cheaper and more predictable than a full build because its shape is known in advance — it is a bounded, contract-defined agent, not an open-ended decomposition.

Mis-sizing is the expensive failure mode. Routing "what's the weather" to a cultivation burns real money; routing "build me a CRM" to a single capability produces garbage. Triage is where that risk is concentrated and controlled — but note that *picking the right capability/agent* depends on resolving the request against cap-reg, which is itself a non-trivial reliability problem (see §6.7).

3.3 The bus (the reusable seam)

A peer-to-peer message fabric. No central node. Every box in this paper speaks to it the same way. The nutrient envelope is the real contract to freeze:

```
{ intent, origin_channel, identity, budget_cap, trace_id, deadline, reply_to }
```

Everything downstream — routing, observability, budget enforcement, the return path — is derived from this envelope.

3.4 Layer 3 — Capability biomes (the 789 + RAGs, via ~4 adapters)

This is where the reusability claim lives — stated precisely (see §1). We write one **transport adapter** per protocol class, and pair it with a **per-capability interface descriptor** held in cap-reg:

- **HTTP-service adapter** — speaks HTTP to any Frank service (research, recon, comp-intel, ...). The *which endpoint, what request shape* comes from that service's cap-reg descriptor.
- **RAG / archive adapter** — speaks the query protocol to research-archive (8975), knowledge-router (9003), topic RAGs.
- **MCP adapter** — speaks MCP to any MCP tool (Gmail, Drive, Brave, Shodan, ...).
- **cap-reg** — the registry (9172) holds the **interface descriptors** (endpoint, request schema, result shape) *and* is the menu triage resolves needs against. Biomes self-announce; descriptors are authored once per capability.

~4 transport adapters + one declarative descriptor per capability — not 789 hand-written integrations. The transport cost is fixed and small; the interface cost is real but declarative and lives in the registry. This is the concrete, honest meaning of "think bigger / what's the reusability."

3.5 Layer 4 — Mycelium orchestration (runs *above* the bus, private by default)

Mycelium owns all composition. Whenever a request needs more than one capability combined, the orchestration is owned by Mycelium — never by triage. Mycelium runs *above* the bus as a client of it, owns its own fan-out, parallelism, and CommitQueue, and is **not** automatically a biome.

The CLI is the *engine's interface, not a human-in-the-loop requirement*. The build cycle (plant → freeze → cultivate → harvest) is executed identically whether the operator invokes it by hand at a terminal **or** a routed request auto-triggers it programmatically. The operator's terminal keeps the full manual control surface (dry-run, concurrency, spend-cap, skip/resume); the product gets the same surface with the gates wired to the front door instead of a prompt. Auto-triggered invocation is made safe by the spend-gate (before a costly build runs) and the promotion gate (before output joins the organism) — not by requiring a person to type the command.

Mycelium serves the two composing tiers from §3.2:

- **Tier 2 — contract-bound agent (bounded composition).** A fixed-shape Mycelium agent defined ahead of time via a NUTRIENTS contract: a known set of capability calls combined in a known way. No goal decomposition — the shape is authored, not discovered. This is the cheap, predictable middle weight triage routes most multi-capability requests to. **v1 ships only pre-**

defined contract-bound agents (testable, cost-predictable); dynamic on-the-fly assembly is deferred to a later capability (§6.2), mirroring the manual→policy promotion philosophy.

- **Tier 3 — cultivation (open-ended composition).** A novel goal decomposed from scratch into HYPHAE and leaves. Gated by the spend-gate.

In both tiers, the **leaves** (or agent steps) reach *back up* to capability-biomes on the bus as their toolset. This is the moment composition and the fleet finally meet: a leaf doing real work can call comp-intel, a RAG, or Gmail through the same bus everything else uses. Because Mycelium runs above the bus, its internals stay private and the CLI-level control we depend on is preserved intact.

The rationale for *not* making every cultivation a biome: a user could build nonsense. We do not want arbitrary user-built artifacts auto-joining the organism. They stay private until they earn promotion. (Note: a pre-defined Tier-2 contract-bound agent is the *opposite* case — it is authored and trusted up front, so it can be a first-class, promotable unit from day one.)

3.6 Layer 5 — Promotion (gated path: cultivation → permanent biome)

This is how the organism grows itself over time without becoming a junk drawer:

```
Cultivation output → Approval gate → Wrapped as adapter → Permanent capability-biome
(works, but        (useful? safe?    (gets a bus adapter    (joins Layer 3; future
private)           non-nonsense?)  + cap-reg entry)        triage routes to it d
```

- **Start: manual approval.** The operator reviews each candidate: is it useful, safe, and worth keeping?
- **Later: policy auto-approve.** Rules auto-promote safe categories; the rest still escalate to manual.

The payoff: every approved promotion makes future triage smarter. Instead of rebuilding a capability via cultivation each time, triage can route directly to the now-permanent biome. The organism accretes proven capability and never accretes junk.

3.7 Cross-cutting concerns (span every layer)

Four concerns thread through the whole system. Notably, because they are uniform and carry the `trace_id`, **they are the natural data source for an operator dashboard** — one coherent story per request, from door to reply.

- **Memory / identity** — who is the caller, and what does Maui remember about them and their org? The colleague↔product boundary lives here (deep-dive deferred).
- **Budget ledger** — per-trace cap drawn from the envelope. Triage reads it; cultivations and capability calls decrement it.
- **Observability** — `trace_id` threads the entire journey: door → triage → bus → biomes → reply. This is what the dashboard renders.

- **Safety tier** — a product caller is not the operator. This gates which capability-biomes are even *visible* to a given tier (e.g. OSINT/recon hidden from the product tier). This is the home of the colleague↔product separation; a dedicated deep-dive is deferred.
-

4. Runtime view — one request, end to end


The structural view shows where things live. The runtime view shows *what happens, in order*.

Worked example: "**Find me 3 suppliers for X and draft outreach.**"

Runtime flow — one request, end to end

The TEMPORAL path (what happens, in order) — complement to the structural diagram. Follows a single intent from iMessage all the way to the capability call and back home.

Front door Triage Tier-2 agent (Mycellium) Bus Adapter Capability (beyond the adapter) Return path



#	Lane	What happens	Wire
1	Front door	iMessage handler wraps the message as an intent nutrient; stamps identity, reply channel, fresh trace_id, and a budget_cap from the caller's tier.	POST bus ← { intent, identity, reply_to:"msg:+1305...", trace_id:7f3a, budget_cap:\$5 }

#	Lane	What happens	Wire
2	Triage	Sizes it. Needs 2 capabilities (supplier search + draft) combined, but not a novel build → Tier 2, composed task . Triage does not compose — it matches the request to a pre-defined contract-bound agent and routes to it. Under cap → no spend-gate.	cap-reg match → agent:"supplier-outreach" (contract: comp-intel → draft)
3	Tier-2 agent (Mycelium)	The contract-bound supplier-outreach agent owns the composition. Per its NUTRIENTS contract it issues two request nutrients to the bus and will combine their results. trace_id rides along.	agent emits: nutrient[supplier-search] → comp-intel · nutrient[draft] → draft (parallel)
4	Adapter	The HTTP-service adapter translates a bus nutrient into the call the real service expects, using comp-intel's interface descriptor from cap-reg (endpoint + request schema). The adapter knows <i>transport</i> ; cap-reg supplies the <i>interface</i> .	descriptor(comp-intel) ⇒ POST http://127.0.0.1:9118/company { name:"X" }
5	Capability (beyond the adapter)	comp-intel runs its own internal pipeline — fans out to Glassdoor/Levels/Blind, uses its fallback chain if a source is down, maybe enriches from a RAG. The bus and adapter never see any of this. The capability is a black box: request in, result out.	comp-intel internal: source A → source B (fallback) → dedupe → RAG enrich → result
6	Return	Adapter wraps the native response as a result nutrient (same trace_id). Bus delivers it back to the Tier-2 agent. Budget ledger decrements actual spend. On error/timeout → an <i>error nutrient</i> instead, and the agent (per contract) decides: retry, fallback, or surface the failure.	{ result, trace_id:7f3a, cost:\$0.12 } → bus → agent
7	Return	The Tier-2 agent waits for both sub-results, composes one coherent answer per its contract, and hands it back to triage, which replies via reply_to . The user gets a single reply on the channel they messaged from. The whole journey is one trace on the dashboard.	imsg send +1305... "Here are 3 suppliers + a draft: ..." (trace 7f3a closed)

Branch — if triage had sized this as Tier 3 (build-anything): it would estimate cost, fire the **spend-gate** back over the front door ("build X, ~\$Y, go?"), and only on approval spawn a Mycelium cultivation that decomposes the goal from scratch. Steps 4–6 then repeat per leaf, with leaves calling capability-biomes exactly as shown.

The key insight of the runtime view: **the transport adapter is a thin translation boundary driven by cap-reg's interface descriptor; the capability behind it is an opaque pipeline of arbitrary depth.** The adapter supplies *transport*; the descriptor supplies the *interface* (endpoint, request schema); together they make the call. The bus does not care how comp-intel finds suppliers — it only cares that comp-intel takes a request nutrient and returns a result nutrient. This encapsulation is what lets transport stay fixed (~4 adapters) even though each capability still needs its own descriptor.

5. Why this is the reusable move

- **Transport cost collapses from $O(N)$ to $O(\text{protocols})$; interface cost stays $O(N)$ but goes declarative.** Adding the 790th capability needs no new integration *code* if it speaks a protocol we already adapt — but it still needs its interface descriptor authored in cap-reg. The win is that per-capability work shrinks from a bespoke integration module to one registry entry. Adding the second front door is free (one more biome). Adding a new orchestration mode reuses the same capability layer.
 - **Mycelium and the fleet finally compose.** Today they are siloed: cultivations build apps; the fleet answers research. With the bus, a cultivation leaf can call any capability, and a capability result can feed a cultivation. They become one system.
 - **Observability is uniform and free.** Because everything carries `trace_id` through one envelope, the dashboard is a read over the bus, not a bespoke integration per surface.
 - **The organism grows safely.** The promotion gate means proven work accretes as permanent capability while junk stays quarantined as private cultivation output.
 - **Front doors are interchangeable.** iMessage today, web/email/API later — each is one biome. The product is not "an iMessage bot"; it is a bus with a swappable mouth.
-

6. Open seams (explicitly deferred, not forgotten)

These are real and need their own treatment before or during build. Listed here so they are not silently dropped.

1. **Failure / timeout paths.** *Resolved (interface)* — see *[failure handoff contract](./mycelium-failure-handoff-contract.md)*. Three handling levels (capability-internal fallback, unchanged · transport/bus · semantic/orchestrator); bus enforces the envelope `deadline` (authoritative) while the capability self-budgets within it; five-class failure taxonomy (`exhausted` / `deadline_hit` / `unreachable` / `bad_request` / `degraded`) carried on a full `ErrorNutrient` schema. ****Recovery policy deferred**** — fallback-biome *selection* waits on §6.7 (cap-reg resolution); partial-result *composition* waits on §6.2 (Tier-2 mechanics). The interface those policies consume is locked.

2. **Tier 2 — pre-defined → dynamic composed tasks.** *Resolved in structure:* multi-capability requests are Tier-2 composed tasks owned by a contract-bound Mycelium agent, not triage and not (necessarily) a cultivation (§3.2, §3.5). **v1 ships only pre-defined contract-bound agents.** The *remaining* open question is narrower: when does a request justify *dynamic* on-the-fly composition (no pre-authored contract), and where is the line between a dynamic Tier-2 assembly and a true Tier-3 cultivation? Dynamic assembly is deferred until the pre-defined pattern is proven, mirroring manual→policy promotion.
3. **Async / long-running replies.** *Resolved (interface) — see [async reply contract](./mycelium-async-reply-contract.md).* One reply path; sync is the degenerate sub-second case. The envelope's `trace_id → reply_to → origin_channel → identity` tuple routes every reply home. Channel is **progress-capable** via a `kind` field (`progress / confirm_required / result`) — the spend-gate is just the first `confirm_required` nutrient, not a separate mechanism; only the terminal `result` is authoritative and settles the ledger. New `relates_to` field added for inbound correlation. **Persistence + correlation policy deferred** to §6.4: this contract requires the routing tuple be *resolvable at delivery*, but the durable store that survives a restart — and the interrupt/queue/ignore policy for a `relates_to` intent — is statefulness.
4. **Statefulness / follow-ups.** "Make it shorter" referring to the previous answer. Does the intent nutrient carry conversation context? Where does conversational state live — front door, memory layer, or trace? **Now the convergence point: both the failure contract (§6.1) and the async-reply contract (§6.3) hand off here.** It owns the durable `trace_id → reply_to → identity` store that lets a reply survive a process restart, and the inbound-correlation policy (interrupt/queue/ignore/fold) for a `relates_to` intent. Next deployment-blocking seam to design.
5. **Safety tier deep-dive.** The colleague↔product boundary. Which capability-biomes are visible to which tier, how identity maps to tier, and how the operator's full-access door coexists with a constrained product door. (Acknowledged as the right home in §3.7; deep-dive deferred by decision.)
6. **The "cultivation as biome" question, revisited.** Resolved *for now* as "no, private by default, promotion-gated." If recursion/observability pressure later argues for cultivations-as-biomes, the promotion gate is the natural place that decision would flip — a promoted cultivation could itself register as a recurring biome.
7. **Intent → capability resolution is itself a reliability problem (load-bearing assumption).** The whole design assumes triage can reliably resolve a natural-language need to the right capability/agent via cap-reg (e.g. `need="supplier-search" → comp-intel`). In this environment that is *not* a solved lookup: the capability-discovery protocol observed that ~50% of capabilities go undiscovered when the cap-reg step is skipped, and keyword enrichment alone only improved discovery by ~1.4pp. Routing quality — not transport — is where a large share of real-world mis-behavior will live. This must be designed (and measured) explicitly: how needs map to descriptors, how ambiguity is handled, and what the fallback is when resolution is low-confidence. Named here as a load-bearing assumption, not an incidental detail.
8. **The spend-gate gates cost, not harm.** Auto-triggered headless invocation (§3.5) is made *cost-safe* by the spend-gate, but a cultivation can do non-cost damage (wrong, misleading, or

destructive output) before the promotion gate ever sees it. A headless-autonomy story needs a harm/safety check distinct from the cost check — what a leaf is allowed to *do* (write where? send to whom?), not just what it may *spend*. Logged here; design deferred.

7. Decisions locked in this session

- **Bus is the universal seam** (forward-compatible; bus concept integrated later in build sequence, design now).
 - **The reusability is ~4 transport adapters + a per-capability interface descriptor in cap-reg** — not "789 integrations become free." Interface cost moves to the registry; it does not vanish.
 - **Names are placeholders** — the seams are load-bearing, the biological vocabulary is not.
 - **Triage sizes and routes; it never composes.** Three tiers: Lookup (single) · Composed task (Tier 2) · Build-anything (Tier 3).
 - **All composition is owned by Mycelium**, above the bus, private by default. A cultivation is not automatically a biome.
 - **The CLI is the engine's interface, not a human-in-the-loop requirement.** The build cycle is invoked identically by an operator typing `mycelium ...` or by a routed request triggering it programmatically. Auto-triggered headless invocation is made safe by the spend-gate + promotion gate.
 - **Tier 2 = contract-bound Mycelium agent.** Pre-defined first (v1), dynamic on-the-fly assembly later — mirroring manual→policy promotion.
 - **Promotion is gated: manual approval to start, policy auto-approve later.**
 - **Cross-cutting concerns double as the dashboard data source.**
 - **Safety tier is the colleague↔product boundary** (deep-dive deferred).
 - **Intent→capability resolution via cap-reg is a load-bearing reliability assumption**, not a given — must be designed and measured (§6.7).
 - **First front door is iMessage + phone**, treated as one biome among future many.
-

8. Suggested next steps (not yet decided)

1. Pick **one** open seam from §6 to design next — failure/timeout is the strongest candidate because it blocks any real deployment.
2. Decide the **build sequence**: does the bus land first, or do we ship the iMessage→trriage→single-capability happy path against the existing fleet *without* the full bus, then retrofit the bus? (Forward-compatible design makes the latter viable.)
3. Decide whether the first build is **dogfooded as a Mycelium cultivation** (per the established operator-authored-hyphae pattern) or implemented directly.

This document is a design artifact, not an implementation plan. The next step in the workflow is to select a scope and produce an implementation plan via the writing-plans process.